

OOLONG: EVALUATING LONG CONTEXT REASONING AND AGGREGATION CAPABILITIES

Anonymous authors

Paper under double-blind review

ABSTRACT

As model context lengths continue to grow, concerns about whether models effectively use the full context length have persisted. While several carefully designed long-context evaluations have recently been released, these evaluations tend to rely on retrieval from one or more sections of the context, which allows nearly all of the context tokens to be disregarded as noise. This represents only one type of task that might be performed with long context. We introduce OOLONG, a benchmark of long-context reasoning tasks that require analyzing individual chunks of text on an atomic level, and then aggregating these analyses to answer distributional questions. OOLONG is separated into two task sets: OOLONG-synth, a set of naturalistic synthetic tasks, where we can easily ablate components of the reasoning problem; and OOLONG-real, a downstream setting which requires reasoning over real-world conversational data. OOLONG requires models to reason over large quantities of examples, to perform both classification and counting in-context, and to reason over temporal and user relations. Even frontier models struggle on OOLONG, with the best model, GPT-5, achieving less than 50% accuracy on both splits at 128K. We release the benchmark examples, code to construct additional evaluation examples for OOLONG-synth, and full outputs and task-specific evaluation results for all models tested to enable further development of models that can reason over large quantities of text. The data and evaluation harness for OOLONG will be released to the research community to build on.

1 INTRODUCTION

In the last several years, the exponentially increasing context lengths of LLMs have enabled many new applications, including reasoning models (Guo et al., 2025), many-shot prompting (Bertsch et al., 2024), and repository-level code generation (Jimenez et al., 2023). However, despite the rapid improvement in this area, measuring effective use of long context windows for complex reasoning tasks remains a challenge. The majority of evaluation datasets involve only simplistic needle-in-a-haystack or retrieval challenges. The few tasks that require *aggregation* over the input, such as text summarization, require evaluation over long-form model responses. Such evaluations are subjective and often expensive to include in an evaluation benchmark. A recent line of work, including tasks that require counting frequent words (RULER; Hsieh et al. (2024)), or identifying the N th occurrence of a particular needle (MRCR; Vodrahalli et al. (2024); OpenAI (2025)), are useful but insufficient measures of information aggregation, only distantly related to real aggregation tasks.¹

In this work, we focus on information aggregation over long inputs. We break down the information aggregation task into a multi-step problem. Each individual subtask is easy to perform for both humans and LLMs. The key challenge is to break down the task into individual subtasks and to aggregate the results of these subtasks. For example, consider the two information aggregation tasks presented in Figure 1. The LM must identify parts of the input relevant to the question, perform a subtask at each relevant part (e.g. a simple classification) and then pool the results from all the subtasks to generate the final answer. A strong long-context reasoning model should be able to do this task in a single pass. Such abilities are useful in many real-world tasks, such as identifying trends in a corpus, building a timeline of medical treatment for a patient, or tracking state in a long

¹For more discussion of current long-context benchmarks and their limitations, see § 5.

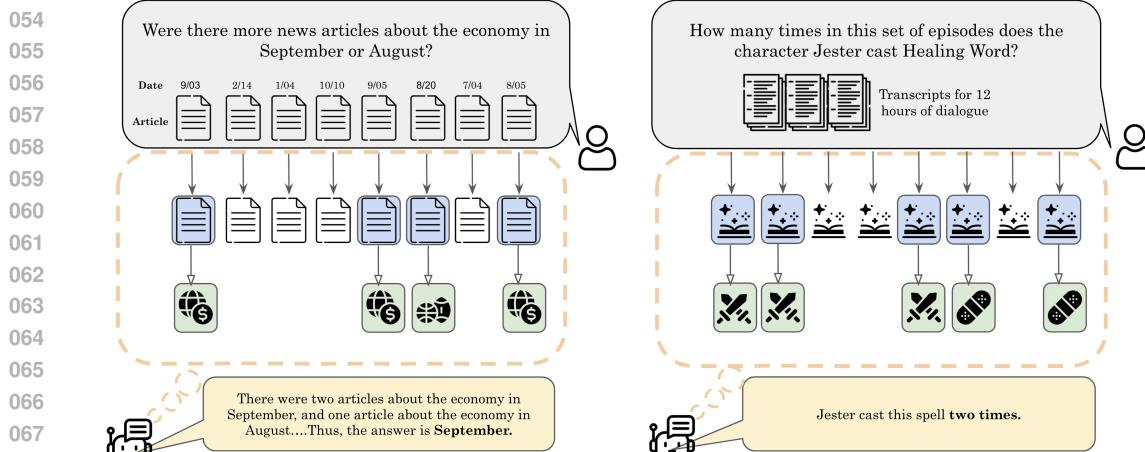


Figure 1: OOLONG poses questions that require performing a multi-step information aggregation process to determine the solution. OOLONG-synth uses ICL-based tasks, which could be easily decomposed and solved iteratively, as a proxy for real-world aggregation tasks over long inputs. OOLONG-real poses challenging information aggregation questions over transcripts from live-action Dungeons & Dragons shows, which can not be easily decomposed into component pieces.

narrative. However, to the best of our knowledge, there are no benchmarks that evaluate LM’s ability to perform information aggregation at scale.

We propose OOLONG, a benchmark that requires multi-hop reasoning over long inputs to produce easily verifiable outputs. OOLONG is separated into two task sets. OOLONG-synth (§2) is a set of naturalistic synthetic tasks constructed from existing in-context learning datasets. These tasks require implicitly labeling the examples in-context to reason over distributional properties of the labels (*counting tasks*), over user-specific patterns (*user tasks*), and over changes in the data over time (*temporal tasks*). OOLONG-real (§3) poses the same types of questions over real data that is not so easily separable into component parts— asking challenging questions about character states and campaign statistics from live-play Dungeons & Dragons roleplaying transcripts, using human-annotated gold answers.

Each task requires identifying the relevant segments of the input; this ranges from only a handful of instances or lines of dialogue to questions that require the use of *every* line of the input. The relevant segments must be classified or categorized, and these individual decisions must be aggregated to produce a final answer. By framing the benchmark around problems that are *simple individually*, we ensure that we are measuring capability on long context reasoning and not accuracy on the underlying task. By requiring identification of relevant context, classification decisions over that context, and numeracy skills to produce statistics about the input, OOLONG requires models to perform multi-step reasoning over long inputs.

Although these component capabilities have been well studied and the frontier models demonstrate strong performance in each (Yen et al. (2024), Agarwal et al. (2024), *inter alia*), we find that models struggle with information aggregation as the input length grows. None of the models benchmarked score higher than 50% on OOLONG at 128K context. We study why this task is so challenging using the more flexible OOLONG-synth, ablating settings that decrease context length dramatically and remove the line-by-line classification task, and find that identification and aggregation of information is the bottleneck, not labeling (§4.1). We believe OOLONG is a usefully challenging evaluation of long-context reasoning abilities.

2 OOLONG-SYNTH

We aim to construct a synthetic information aggregation task that allows fine-grained control over the types of information and number of steps necessary to solve the task. To do this, we need building blocks: simple, short-context documents to build into a longer collection. We construct OOLONG-

Dataset	Task	# Labels	Input Len
Spam (Almeida et al., 2011)	SMS spam classification	2	57
TREC-QC-coarse (Li & Roth, 2002; Hovy et al., 2001)	Question type classification	6	39
AGNews (Zhang et al., 2015)	Headline topic classification	4	90
App Reviews (Zur, 2017)	Review sentiment classification	2 [†]	49
Pavlick Formality (Lahiri, 2015; Pavlick & Tetreault, 2016)	Formality classification	2 [†]	51
IMDB reviews (Maas et al., 2011)	Sentiment analysis	2 [†]	376
HiTZ Negation (García-Ferrero et al., 2023)	Verify claims about definitions	2	45
Yahoo Topics (Zhang et al., 2015)	Question topic classification	10	74
MultiNLI (Williams et al., 2018)	Entailment	3	70
Metaphors (Bizzoni & Lappin, 2018)	Metaphor meaning validation	2	51

Table 1: Datasets used for OOLONG-synth. [†] indicates cases where we combined similar labels to produce a smaller label-set for OOLONG. The input length is the average per-instance length (in Llama 2 tokens), including our added date and user metadata.

synth by constructing challenging corpus-level questions over existing in-context learning (ICL) datasets.

2.1 DATA

We collect 10 common text classification datasets with between 2 and 10 labels. We select for tasks that are possible for the authors to perform without difficulty (and validate that these tasks are similarly simple for models during our dataset filtering). Table 1 describes the datasets in more detail. We split the data into two validation tasks and eight test tasks. The validation tasks are selected to avoid underlying task overlap with test tasks (in the style of Min et al. (2022)), and so that there is both a 2-label and a multi-label (6-label) validation task. For the sentiment and formality datasets with more than two labels, we reduce the label space to a single positive and a single negative label, to reduce the difficulty of the task.

Data filtering In a normal in-context learning task, if a few examples are unusually hard or even mislabeled, this has a small impact on the overall score. Because we intend to require aggregation across many ICL examples at once, one particularly difficult example can affect the score on many questions downstream. However, many ICL datasets are known to contain mislabelings (Ying & Thomas, 2022; Chong et al., 2022; Klie et al., 2023).

We perform an additional screening step to remove mislabeled or unusually hard examples. We select two models that we do not expect to be substantially stronger than the models we are evaluating: GPT-4.1 nano and Llama 4 Maverick. We perform zero-shot ICL with a minimal instruction that provides the label space and type of task. We then exclude all examples that *both* models get incorrect. We report the percentage of examples screened out by dataset in Table 1, and Appendix A analyzes excluded and validated datapoints for each dataset. Many of the excluded instances are clearly mislabeled; the remainder generally represent challenging cases.

2.2 CONTEXT WINDOW CONSTRUCTION

We determine the number of examples in the context window by using an estimate of the number of tokens required for each per example, for the general task instructions, and for the specific question. Following Yen et al. (2024), we compute these averages using the Llama 2 tokenizer, and use approximately 500K tokens of context for each estimate. We then estimate the number of examples to use by determining how many average-length examples would reach 95% of the target context length.

Dataset	% Removed
Spam	0.635%
TREC-Q-coarse	0.048%
AGNews	0.026%
App Reviews	0.051%
Formality	0.108%
IMDB	0.042%
Metaphors	0.000%
MultiNLI	0.109%
Negation	0.016%
Yahoo Topics	0.188%

Table 2: Only a small fraction of examples are removed during validation for each source dataset.

We sample a distribution over the label classes so that the model cannot use any information about the true distribution over labels (e.g., by guessing that most sentiment classification datasets are balanced). Then, we sample examples to fill this distribution; if there are less total examples than required, we sample with replacement. For each instance, we sample a date and user ID. User IDs are drawn such that 80% of instances have an ID in the 20% of IDs that are most common. Dates are drawn uniformly with replacement from an approximately 40-month range.

Once a context window is constructed, we use it for 25 questions; reusing the same context window allows for prompt caching, reducing the time to evaluate local models and the cost to evaluate on most APIs. At each context length, we sample two context windows per dataset, for a total of 50 questions per dataset or 400 total questions in the test split per context length. We construct OOLONG-synth questions for every power of 2 from 1K to 4M.

Questions We construct three types of OOLONG-synth questions, in order of increasing complexity. *Counting* questions concern simple statistical properties of the label distribution, ranging from identifying the most frequent label to determining the number of examples with each label class. If the label distribution was provided, all tasks in this set should be trivial; these tasks can be seen implicitly as the task of labeling every example and reporting summary statistics. *User information* questions require additional cross-reference with the *user ID* field. *Timeline* questions ask about changes in distribution before or after a certain date, between years, or between months across years. This is a more challenging than user information questions because it requires reasoning about *before* or *after* a date, rather than matching a list of IDs (see §B.1 in Appendix for questions).

Instructions For each dataset, we provide a brief (one-sentence) description of the task and label space, as well as the number of examples in-context. We provide instructions at the start and end of the input, as suggested by OpenAI’s long context prompting guide.² However, we provide the question only at the end, to enable prompt caching.

2.3 EVALUATION

Baseline A random baseline is non-trivial for these tasks, because the output space for each question varies. We construct an algorithm for a random baseline with the following rules: (1) in cases where there is an n -way choice (e.g. choosing between n labels or choosing a month where some criteria occurs), we select an answer from the set of valid answers at random; (2) if the question requires a numerical answer, we return $N/|L|$, where N is the number of data points in context and L is the labelset; (3) if the question requires a date or user ID, we sample this from the list of dates/IDs in context at random. We compute the random baseline performance as the expected value of this procedure over the dataset.

Parsing answers The task prompts specify an output format. We make a best-effort attempt to parse answers using this format; if this template is not present in the output, we choose the end of the input as the candidate answer and mark this as a low-confidence parse. Generally, this occurs if the model runs out of output tokens before providing an answer.

Scoring For questions that require a label, date, user ID, or comparison (e.g. “greater than” or “less than”), we score on exact match. For questions that require a numerical answer, we set the score to be

$$\text{score}(\hat{y}) = 0.75^{|y - \hat{y}|}$$

This allows for partial credit for answers close to the exact value.

3 OOLONG-REAL

Although OOLONG-synth allows for a high degree of control over the types of input, these inputs are composed of independent examples from the source dataset. An iterative setup could complete these

²https://cookbook.openai.com/examples/gpt4-1_prompting_guide#prompt-organization

216	<i>Counting</i>
217	Total number of rolls in this episode?
218	What is the count of Crits? (natural rolls of value 1 or 20)?
219	How many {spell type} spells were cast during this episode?
220	How many characters cast {spell name} spell all across episodes?
221	What is the second spell cast in the episode {episode index}?
222	
223	<i>Enumeration</i>
224	What are the first {count} spells cast in this episode? Return a comma separated list.
225	List the last spell cast in each episode? Return a comma separated list.
226	
227	<i>Indexing</i>
228	What is the cumulative total of rolls by the end of episode {episode index}? Count the number of rolls and not the values of the rolls.
229	What is the second spell cast in the episode {episode index}?
230	List the last spell cast in each episode? Return a comma separated list.
231	

Table 3: Question types covered in OOLONG-real dataset.

tasks by calling a model to label each individual example and then deterministically aggregating the resulting labels. We note that this is true for many capability evaluation tasks; for instance, Needle-in-a-Haystack tasks could be solved almost trivially by asking a model if each of the N input sentences individually contains the needle. Although we believe tasks like these still have merit as capability evaluations, OOLONG-synth does not answer the question of why information aggregation is a useful capability. For this, we must turn to a more realistic setting.

We complement OOLONG-synth with questions derived from real conversational data. OOLONG-real is compiled from the transcripts of a Dungeons and Dragons (D&D) role-playing game, where a group of players collaboratively build a story through in-character actions and success depends on rolling dice. Stories unfold over narratives (“campaigns”) that span dozens to hundreds of episodes, with each episode generally lasting 4-5 hours of play. Compiling aggregate statistics from real conversational data is *not* simple to reframe as an iterative task. These transcripts involve several levels of conversion, from out-of-character discussion to rules discussion to in-character actions and speech. Though lightly edited for readability, they reflect naturalistic speech instead of carefully planned written text. Because the conversation is unscripted and can involve tangents or side channels, conversational turns require variable amounts of prior context to resolve. In some instances, the same event (e.g. casting a spell) is discussed for many turns, or brought up again after a long interlude; in other cases, a prior event is “retconned” or revised post hoc.

In Dungeons and Dragons, there are limitations to how frequently characters can take certain actions (e.g., cast certain types of spells). Additionally, fans are often interested in metadata of the play (e.g., whether a certain person is particularly unlucky with their dice rolls or if a character uses a signature spell more or less as the campaign progresses). Because these shows are extremely long and wide-ranging and improvised live, the creators do not plan for or provide this type of information. In lieu of an official source, this information is often annotated, meticulously, and with multiple levels of verification, by dedicated fans of the work.

We take this as an example of an information aggregation task in the wild. We consider the series Critical Role and the data compiled by CritRoleStats,³ a fan project that tracked per-episode statistics for the first several campaigns of Critical Role. We devise questions related to characters, dice rolls, and spells cast during episodes and use human-labeled CritRoleStats to compute gold answers.

3.1 DATASET COMPILATION

For OOLONG-real, we consider two campaigns from the Critical Role TV series. We used episode transcripts from the Critical Role Dungeons and Dragons Dataset (CRD3) (Rameshkumar & Bailey, 2020), which includes full episode transcripts from the first two campaigns of the Critical Role TV

³<https://www.critrolestats.com/>

270 series. For our testbed, we used the first campaign, which consists of 115 episodes. Each line in the
 271 transcript includes an utterance with the player name labeled.
 272

273 **QA pairs** We use the game statistics compiled by the authors of CritRoleStats. Specifically, we
 274 utilize statistics about dice rolls and spells cast in each episode. We design a set of questions that
 275 cover a variety of information aggregation tasks; see Table 3 for some examples. We include ques-
 276 tions that require processing of single- or multiple-episode transcripts. For multi-episode questions,
 277 we concatenate transcripts and use delimiters to highlight the start and end of each transcript. See
 278 Table 18 and Table 19 of the Appendix for a full list of OOLONG-real questions. We include varying
 279 context windows to evaluate the model’s ability to use long-context reasoning and aggregation ca-
 280 pabilities. Using a single episode transcript as a context unit, we include context windows ranging
 281 from one to 24 episode transcripts. This covers input lengths of 55K to 1.3M tokens.

282 3.2 EVALUATION

283 **Parsing** Our task prompt requires the model to place the answer in $\boxed{\{ \}}$. If the answer cannot
 284 be extracted successfully, we attempt to extract the answer for a given question using GPT-5-nano.
 285

286 **Scoring** OOLONG-real contains three types of answers: numeric, string, and a list of strings. For
 287 numerical answers, we use the same scoring scheme as OOLONG-synth. We use exact match for
 288 string answer types and set overlap for answers of a type list.
 289

291 4 RESULTS AND ANALYSIS

294 On OOLONG, we benchmark a strong suite of frontier models to study their information aggregation
 295 capabilities over long context. We include GPT-5, GPT o-series, Claude, DeepSeek R1, and Llama
 296 4. We include a mix of models of varying sizes, levels of reasoning, and long-context capabilities.
 297

298 4.1 OOLONG – LEADERBOARD

300 In Table 4, we report the results for both OOLONG-synth and OOLONG-real. Following prior long-
 301 context benchmarks, we measure model performance at increasing context window sizes and report
 302 a task average. For OOLONG-synth, we use the standard context window sizes of up to 128K tokens.
 303 For OOLONG-real, we report scores at increasing number of episode transcripts (up to 3 episodes).
 304 GPT-5 performs the best in both subsets, followed by o3, GPT-5-mini and Claude-Sonnet-4. No-
 305 tably, at 128K context window size all our evaluated models obtain a score under 50. DeepSeek R1
 306 does well on OOLONG-real, but struggles to perform well on OOLONG-synth (see §4.2. Llama-4-
 307 Maverick fails in both subsets.
 308

309 Model	Avg.	OOLONG-synth					OOLONG-real			
		8K	16K	32K	64K	128K	Avg.	55K	118K	175K
GPT-5	67.6	80.6	80.0	71.8	60.0	45.7	47.0	58.7	45.7	36.5
o3	61.1	82.3	76.0	61.6	47.6	37.8	36.7	50.6	33.6	26.0
GPT-5 mini	62.0	80.1	73.7	62.5	51.4	42.2	34.5	49.9	29.9	23.9
Claude Sonnet 4	55.3	70.2	60.3	51.1	48.6	46.2	36.8	50.6	33.0	26.7
o4-mini	54.7	77.8	62.4	51.1	44.4	37.9	27.1	41.7	21.8	17.9
GPT-5 nano	49.1	68.0	53.8	43.2	42.7	38.0	31.0	43.1	26.8	23.2
Deepseek R1 0528	11.5	12.4	12.7	11.6	11.6	9.1	32.0	47.9	27.3	20.8
Llama 4 Maverick	14.3	12.3	14.5	12.9	16.5	15.4	2.1	2.5	2.1	1.6

319 Table 4: OOLONG results on a number of strong models. All models we test support at least 200K
 320 context; thus, we report OOLONG scores as an average over scores on 8K-175K inputs. OOLONG-
 321 synth is a more information-dense task than OOLONG-real; while the best model remains the same,
 322 the ranking of other models shifts slightly between the two settings. Models are sorted by the average
 323 between the two benchmarks.

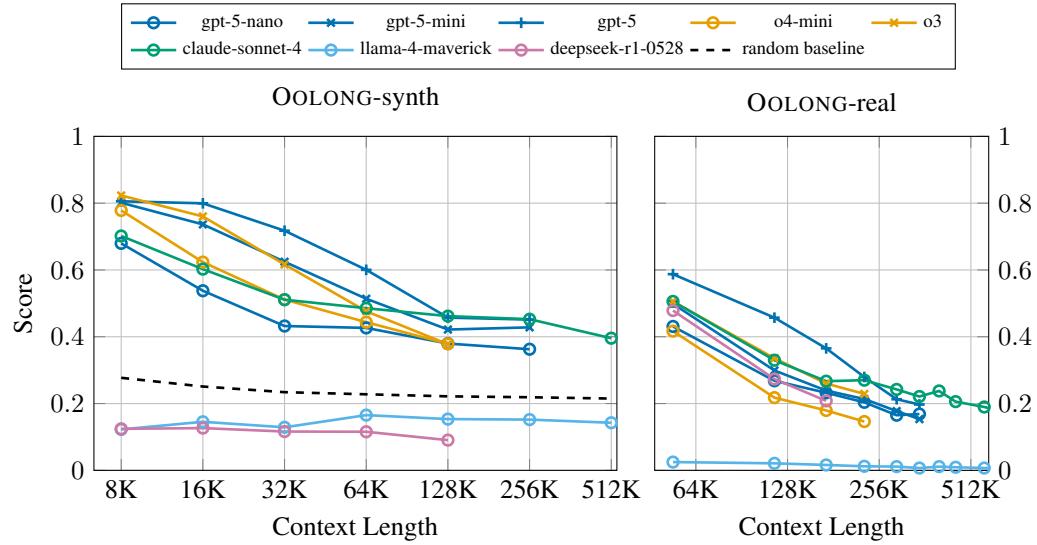


Figure 2: Scores by context window length for OOLONG-synth and OOLONG-real.

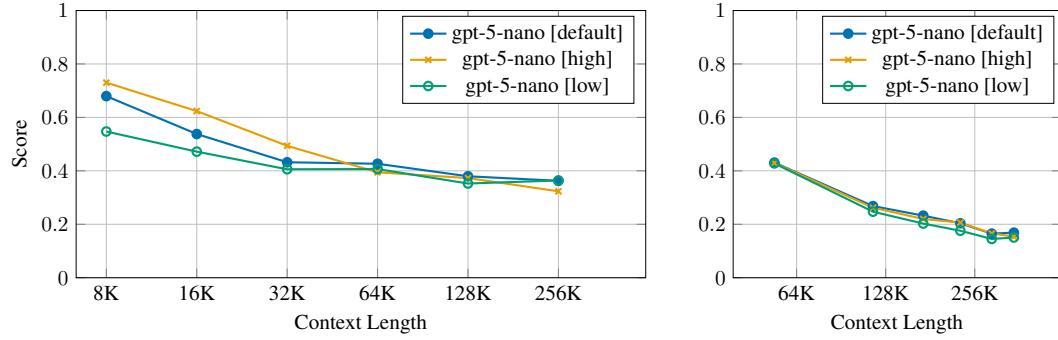


Figure 3: Comparison across reasoning levels.

In Figure 2, we visualize the model performance for context windows up to 512K tokens. As expected, we see a significant drop in performance at higher context windows. The two splits are of similar difficulty at the same context length, although direct comparison is challenging because real data do not necessarily align with fixed context window buckets.⁴

4.2 IMPACT OF REASONING LEVEL

We compare “high” and “low” reasoning effort for GPT-5-nano. Figure 3 shows the results for both the splits of OOLONG. Although OOLONG is a reasoning-intensive task, specifying a higher reasoning effort is only useful for short contexts; after 64k, there is little discernible difference between reasoning levels, with the “high” reasoning setting even slightly underperforming “low” reasoning at 256k on OOLONG-synth. Although it is difficult to draw firm conclusions without the ability to view the reasoning trace, we hypothesize that at context lengths where there is sufficient remaining room in the context window to enumerate labels for each example in-context, adding more reasoning effort may encourage the model to take this strategy. However, since the default routing for this model performs reasonably well at all lengths and especially for longer inputs, we do not explicitly specify a reasoning level for the remainder of the runs.

⁴While performance on OOLONG-real starts lower, this is because the shortest inputs in this dataset are a single episode, with an average length of 55K tokens; models perform similarly on OOLONG-synth around the same context length.

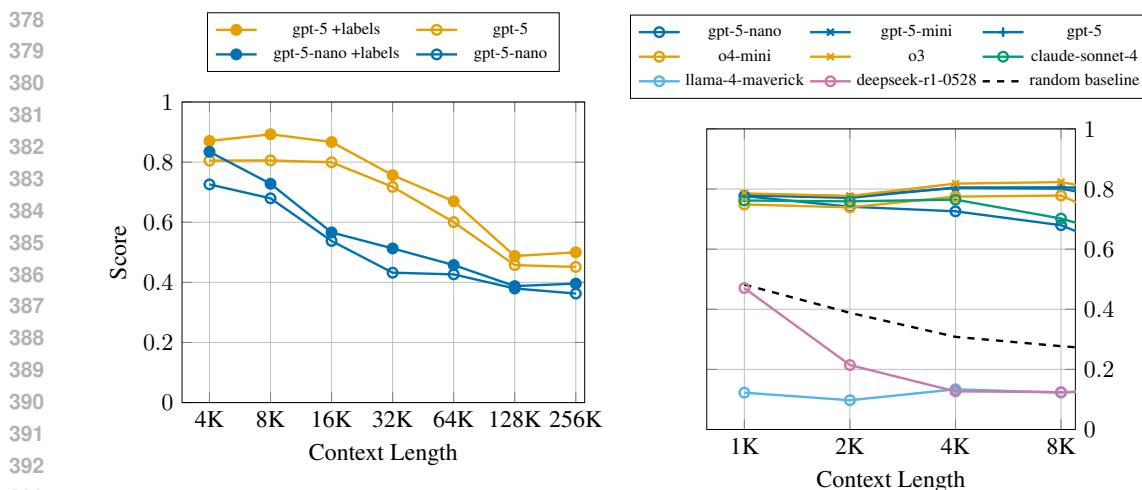


Figure 4: Comparison on OOLONG-synth: (a) we provide the gold labels in the input. This leads to a consistent but small improvement, (b) short context performance; while the top models have similar short-context performance, differences emerge as the context length grows.

4.3 WHAT HAPPENED TO DEEPSEEK R1?

Deepseek R1 is a strong reasoning model and outperforms GPT-5 nano and o4-mini on OOLONG-real. However, on OOLONG-synth, it achieves performance below the random baseline. What causes this discrepancy?

In manual inspection of the traces, we observe two pathologies in the reasoning traces for OOLONG-synth that are not present in traces for OOLONG-real. First, the traces for the synthetic split often end mid-sentence, without outputting an answer at all. We hypothesize that, because OOLONG-synth is such an information-dense task, the model’s apparent strategy of labeling each example before deciding which are relevant results in running out of context tokens. Second, many traces use a substantial number of tokens deciding whether or not the model should label the examples in-context. Appendix C shows example traces from both settings.

4.4 SIMPLER SETTINGS FOR OOLONG-SYNTH

Shorter context Most models show declining performance with context length. We consider a short-context version of OOLONG-synth, with inputs between 1k and 4k tokens, in Figure 4. While models perform better on this task, several still struggle, showing that even short-context aggregation remains challenging. Performance differences between the top models are difficult to distinguish in the short-context regime, suggesting that these models have the ability to perform the task at some context length. However, no model exceeds 85% performance at any context length.

Aggregation without classification For OOLONG-synth, we can construct an easier version of the task by providing the label for each ICL example in-context. This reduces the task for most questions to simply identifying the relevant instances and summing the occurrence of each label type. Figure 4 shows the behavior of GPT-5 and GPT-5-nano with and without labels provided in-context. As expected, adding labels improves accuracy; however, this improvement ranges from 10.9 points to only 0.79 points. We do not see consistently higher gains from providing labels in longer inputs, which suggests that the lower performance at longer context lengths is not primarily due to an accumulation of mislabeling errors. Additionally, the improvement is not larger for GPT-5-nano than it is for the more powerful GPT-5, suggesting that the performance difference observed between these models is not due to differing ability to perform the classification task. This is by design; our classification task validation was designed to eliminate overly challenging or misleading examples, as the ability to aggregate information is the main capability we aim to measure.

432 **5 RELATED WORK**

433

434 **Long-context benchmarks** A variety of long-context benchmarks have included some type of
 435 aggregation or reasoning-focused task. RULER (Hsieh et al., 2024) benchmark includes synthetic
 436 retrieval tasks at varying context lengths. It includes multi-hop tracing and aggregation tasks. HEL-
 437 MET (Yen et al., 2024) expands the tasks in RULER to include tasks related to the downstream use
 438 of LMs (reranking, ICL, LongQA, summarization). LongMemEval (Wu et al., 2025) and multi-
 439 round coreference resolution tasks (MRCR) evaluate long-context capabilities with conversational
 440 data. MRCR was first introduced in Gemini (Vodrahalli et al., 2024) and was further extended in
 441 OpenAI’s MRCR (OpenAI, 2025). Unlike the standard NIAH task, the needles and the distrac-
 442 tors in MRCR are sampled from the same distributions. OpenAI’s MRCR includes variants of 2,
 443 4, and 8 needles, and the task involves retrieving the i th instance of one of the needles. The docu-
 444 ments in MRCR are synthetic conversations generated using GPT-4o. TLDM (Hamilton et al., 2025)
 445 includes two related tasks that query about a character last known location and time passed in a nar-
 446 rative setting, but only considers public-domain novels, which likely suffer from data contamination.
 447 ZeroSCROLLS (Shaham et al., 2023) introduces a task of identifying the percentage of reviews on
 448 Amazon for a product that are positive, which is the closest conceptual ancestor of our work; Oo-
 449 LONG-synth encompasses this type of task along with other aggregation tasks. GSM-infinite (Zhou
 450 et al., 2025) stress tests long-context reasoning through an adaptable framework for varying task
 451 difficulty and length; our work is complementary, as their focus is primarily on increasingly difficult
 452 mathematical reasoning through synthetically generated problems, while we focus on a more noisy,
 453 naturalistic text setting with less challenging mathematical reasoning required. Finally, procedural
 454 generation benchmarks (Ye et al., 2025) measure a different type of reasoning task by evaluating the
 455 ability to plan and produce coherent long outputs.

456 In addition to the long-context benchmarks above, prior work has also shown that LM performance
 457 on downstream tasks deteriorates at longer context lengths. Levy et al. (2024) and Shi et al. (2023)
 458 studied the adverse effects of adding irrelevant context to the model input.

459 **Dungeons & Dragons data** Several prior works consider D&D as a potential testbed for model ca-
 460 pabilities; for instance, we use a version of the Critical Role transcripts preprocessed by Rameshku-
 461 mar & Bailey (2020), who proposed an episode summarization task using fan-written summaries.
 462 D&D data has also been used to study theory of mind (Zhou et al., 2023), user assistant development
 463 Zhu et al. (2023b), and dialogue generation, with game state either inferred from forum-based games
 464 (Callison-Burch et al., 2022) or captured through an external tool (Zhu et al., 2023a); to the best of
 465 our knowledge, we are the first to use fan annotations of gold labels and to consider the generation
 466 of these statistics as a task in its own right.

467 **Aggregation as an NLP concept** Goldman et al. (2024) argue for the definition of long context
 468 tasks in terms of the information dispersion and input scope required to find the answer; under this
 469 taxonomy, OOLONG is high dispersion (because relevant info is distribution over the full context
 470 length) and high scope (because most of the input is necessary for the task, particularly for OOLONG-
 471 synth). DeYoung et al. (2024) measure synthesis in multi-document summarization by asking if
 472 summaries convey the consensus opinion of the input documents; this is a related task to the counting
 473 tasks in OOLONG-synth, but framed as a measure of summarization capabilities instead of long
 474 context capabilities, and thus requires different affordances related to evaluating generated outputs.

475 **6 CONCLUSION**

476

477 We introduce OOLONG, a challenging long-context information aggregation benchmark in two parts.
 478 OOLONG-synth uses synthetic aggregation tasks over ICL data to enable finer-grained control of
 479 the benchmark settings, while OOLONG-real poses questions over real long-context conversational
 480 data and human-annotated labels. On both splits, models struggle, with performance dropping with
 481 increasing context length even when controlling for the potential compounding of mislabeling errors.
 482 We see substantial headroom between strong open source models and API-based models on this task,
 483 particularly OOLONG-synth, where Deepseek-R1 struggles to plan the reasoning quantity. Overall,
 484 our results suggest that there is still a long way to go in designing robust long-context aggregation
 485 capabilities for LLMs.

486 ETHICS STATEMENT
487488 This work is intended to further the evaluation of long-context LLMs. We do not foresee any par-
489 ticular additional risks introduced by our evaluation, although information aggregation, like nearly
490 any LM capability, could be useful for both beneficial and harmful actors.491 CritRoleStats, whose data we use in constructing gold labels for the OOLONG-real split, explicitly
492 allows the use of their data for statistical analyses in their FAQ, so long as they are credited.⁵493 In the course of this work, we used LM assistance in the writing of some data preprocessing, analy-
494 sis, and visualization scripts.496 REPRODUCIBILITY STATEMENT
497498 We release code for dataset construction and evaluation, in addition to the final dataset version. API-
499 based models typically exhibit some small variability across API calls, so exact replication of scores
500 may be infeasible; however, we will release the model outputs for each model run reported in the
501 paper.502 REFERENCES
503

504 Software applications user reviews. 2017.

505 Rishabh Agarwal, Avi Singh, Lei M. Zhang, Bernd Bohnet, Luis Rosias, Stephanie Chan, Biao
506 Zhang, Ankesh Anand, Zaheer Abbas, Azade Nova, John D. Co-Reyes, Eric Chu, Feryal Be-
507 bahani, Aleksandra Faust, and Hugo Larochelle. Many-shot in-context learning, 2024. URL
508 <https://arxiv.org/abs/2404.11018>.512 Tiago A. Almeida, Jose Maria Gomez Hidalgo, and Akebo Yamakami. Contributions to the study
513 of sms spam filtering: New collection and results. In *Proceedings of the 2011 ACM Symposium
514 on Document Engineering (DOCENG'11)*, 2011.515 Amanda Bertsch, Maor Ivgi, Emily Xiao, Uri Alon, Jonathan Berant, Matthew R Gormley, and
516 Graham Neubig. In-context learning with long-context models: An in-depth exploration. *arXiv
517 preprint arXiv:2405.00200*, 2024.519 Yuri Bizzoni and Shalom Lappin. Predicting human metaphor paraphrase judgments with deep
520 neural networks. In Beata Beigman Klebanov, Ekaterina Shutova, Patricia Lichtenstein, Smaranda
521 Muresan, and Chee Wee (eds.), *Proceedings of the Workshop on Figurative Language Processing*,
522 pp. 45–55, New Orleans, Louisiana, June 2018. Association for Computational Linguistics. doi:
523 10.18653/v1/W18-0906. URL <https://aclanthology.org/W18-0906/>.524 Chris Callison-Burch, Gaurav Singh Tomar, Lara Martin, Daphne Ippolito, Suma Bailis, and David
525 Reitter. Dungeons and dragons as a dialog challenge for artificial intelligence. In *Proceedings
526 of the 2022 Conference on Empirical Methods in Natural Language Processing*, pp. 9379–9393.
527 Association for Computational Linguistics, 2022. doi: 10.18653/v1/2022.emnlp-main.637. URL
528 <http://dx.doi.org/10.18653/v1/2022.emnlp-main.637>.530 Derek Chong, Jenny Hong, and Christopher D. Manning. Detecting label errors by using pre-trained
531 language models, 2022. URL <https://arxiv.org/abs/2205.12702>.532 Jay DeYoung, Stephanie C. Martinez, Iain J. Marshall, and Byron C. Wallace. Do multi-document
533 summarization models synthesize? *Transactions of the Association for Computational Linguis-
534 tics*, 12:1043–1062, 2024. doi: 10.1162/tacl_a_00687. URL <https://aclanthology.org/2024.tacl-1.58/>.537 Iker García-Ferrero, Begoña Altuna, Javier Alvez, Itziar Gonzalez-Dios, and German Rigau. This
538 is not a dataset: A large negation benchmark to challenge large language models. In Houda5<https://www.critrolestats.com/faqs>

- 540 Bouamor, Juan Pino, and Kalika Bali (eds.), *Proceedings of the 2023 Conference on Empirical
541 Methods in Natural Language Processing*, pp. 8596–8615, Singapore, December 2023.
542 Association for Computational Linguistics. doi: 10.18653/v1/2023.emnlp-main.531. URL
543 <https://aclanthology.org/2023.emnlp-main.531>.
- 544 Omer Goldman, Alon Jacovi, Aviv Slobodkin, Aviya Maimon, Ido Dagan, and Reut Tsarfaty. Is it
545 really long context if all you need is retrieval? towards genuinely difficult long context NLP. In
546 Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen (eds.), *Proceedings of the 2024 Conference
547 on Empirical Methods in Natural Language Processing*, pp. 16576–16586, Miami, Florida, USA,
548 November 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.emnlp-main.
549 924. URL <https://aclanthology.org/2024.emnlp-main.924>.
- 550 Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu
551 Zhang, Shirong Ma, Xiao Bi, et al. Deepseek-r1 incentivizes reasoning in llms through reinforce-
552 ment learning. *Nature*, 645(8081):633–638, 2025.
- 553 Sil Hamilton, Rebecca M. M. Hicke, Matthew Wilkens, and David Mimno. Too long, didn't model:
554 Decomposing llm long-context understanding with novels, 2025. URL <https://arxiv.org/abs/2505.14925>.
- 555 Eduard Hovy, Laurie Gerber, Ulf Hermjakob, Chin-Yew Lin, and Deepak Ravichandran. To-
556 ward semantics-based answer pinpointing. In *Proceedings of the First International Confer-
557 ence on Human Language Technology Research*, 2001. URL <https://www.aclweb.org/anthology/H01-1069>.
- 558 Cheng-Ping Hsieh, Simeng Sun, Samuel Kriman, Shantanu Acharya, Dima Rekesh, Fei Jia, and
559 Boris Ginsburg. RULER: What's the real context size of your long-context language models? In
560 *First Conference on Language Modeling*, 2024. URL <https://openreview.net/forum?id=kIoBbc76Sy>.
- 561 Carlos E Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik
562 Narasimhan. Swe-bench: Can language models resolve real-world github issues? *arXiv preprint
563 arXiv:2310.06770*, 2023.
- 564 Jan-Christoph Klie, Bonnie Webber, and Iryna Gurevych. Annotation error detection: Analyzing
565 the past and present for a more coherent future. *Computational Linguistics*, 49(1):157–198, 03
566 2023. ISSN 0891-2017. doi: 10.1162/coli_a_00464. URL https://doi.org/10.1162/coli_a_00464.
- 567 Shibamouli Lahiri. SQUINKY! A corpus of sentence-level formality, informativeness, and implica-
568 ture. *arXiv preprint arXiv:1506.02306*, 2015.
- 569 Mosh Levy, Alon Jacoby, and Yoav Goldberg. Same task, more tokens: the impact of input length
570 on the reasoning performance of large language models. In Lun-Wei Ku, Andre Martins, and
571 Vivek Srikumar (eds.), *Proceedings of the 62nd Annual Meeting of the Association for Com-
572 putational Linguistics (Volume 1: Long Papers)*, pp. 15339–15353, Bangkok, Thailand, August
573 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.818. URL
574 <https://aclanthology.org/2024.acl-long.818/>.
- 575 Xin Li and Dan Roth. Learning question classifiers. In *COLING 2002: The 19th Interna-
576 tional Conference on Computational Linguistics*, 2002. URL <https://www.aclweb.org/anthology/C02-1150>.
- 577 Andrew L. Maas, Raymond E. Daly, Peter T. Pham, Dan Huang, Andrew Y. Ng, and Christopher
578 Potts. Learning word vectors for sentiment analysis. In *Proceedings of the 49th Annual Meeting
579 of the Association for Computational Linguistics: Human Language Technologies*, pp. 142–150,
580 Portland, Oregon, USA, June 2011. Association for Computational Linguistics. URL <http://www.aclweb.org/anthology/P11-1015>.
- 581 Sewon Min, Mike Lewis, Luke Zettlemoyer, and Hannaneh Hajishirzi. Metaicl: Learning to learn
582 in context, 2022. URL <https://arxiv.org/abs/2110.15943>.

- 594 OpenAI. OpenAI MRCR: Long context multiple needle in a haystack benchmark, 2025. URL
595 <https://huggingface.co/datasets/openai/mrcr>.
596
- 597 Ellie Pavlick and Joel Tetreault. An empirical analysis of formality in online communication. *Trans-*
598 *actions of the Association for Computational Linguistics*, 2016.
- 599 Revanth Rameshkumar and Peter Bailey. Storytelling with dialogue: A Critical Role Dungeons
600 and Dragons Dataset. In Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel Tetreault (eds.),
601 *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp.
602 5121–5134, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/
603 2020.acl-main.459. URL <https://aclanthology.org/2020.acl-main.459/>.
604
- 605 Uri Shaham, Maor Ivgi, Avia Efrat, Jonathan Berant, and Omer Levy. Zeroscrolls: A zero-shot
606 benchmark for long text understanding, 2023. URL <https://arxiv.org/abs/2305.14196>.
607
- 608 Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H. Chi, Nathanael
609 Schärli, and Denny Zhou. Large language models can be easily distracted by irrelevant context.
610 In Andreas Krause, Emma Brunskill, Kyunghyun Cho, Barbara Engelhardt, Sivan Sabato, and
611 Jonathan Scarlett (eds.), *Proceedings of the 40th International Conference on Machine Learning*,
612 volume 202 of *Proceedings of Machine Learning Research*, pp. 31210–31227. PMLR, 23–29 Jul
613 2023. URL <https://proceedings.mlr.press/v202/shi23a.html>.
614
- 615 Kiran Vodrahalli, Santiago Ontanon, Nilesh Tripuraneni, Kelvin Xu, Sanil Jain, Rakesh Shivanna,
616 Jeffrey Hui, Nishanth Dikkala, Mehran Kazemi, Bahare Fatemi, Rohan Anil, Ethan Dyer, Siamak
617 Shakeri, Roopali Vij, Harsh Mehta, Vinay Ramasesh, Quoc Le, Ed Chi, Yifeng Lu, Orhan Firat,
618 Angeliki Lazaridou, Jean-Baptiste Lespiau, Nithya Attaluri, and Kate Olszewska. Michelangelo:
619 Long context evaluations beyond haystacks via latent structure queries, 2024. URL <https://arxiv.org/abs/2409.12640>.
620
- 621 Adina Williams, Nikita Nangia, and Samuel Bowman. A broad-coverage challenge corpus for sen-
622 tence understanding through inference. In *Proceedings of the 2018 Conference of the North Amer-
623 ican Chapter of the Association for Computational Linguistics: Human Language Technologies,
624 Volume 1 (Long Papers)*, pp. 1112–1122. Association for Computational Linguistics, 2018. URL
625 <http://aclweb.org/anthology/N18-1101>.
626
- 627 Di Wu, Hongwei Wang, Wenhao Yu, Yuwei Zhang, Kai-Wei Chang, and Dong Yu. Longmemeval:
628 Benchmarking chat assistants on long-term interactive memory. In *The Thirteenth International
629 Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=pZiyCaVuti>.
630
- 631 Xi Ye, Fangcong Yin, Yinghui He, Joie Zhang, Howard Yen, Tianyu Gao, Greg Durrett, and Danqi
632 Chen. Longproc: Benchmarking long-context language models on long procedural generation,
633 2025. URL <https://arxiv.org/abs/2501.05414>.
634
- 635 Howard Yen, Tianyu Gao, Minmin Hou, Ke Ding, Daniel Fleischer, Peter Izsak, Moshe Wasserblat,
636 and Danqi Chen. Helmet: How to evaluate long-context language models effectively and thor-
oughly. *arXiv preprint arXiv:2410.02694*, 2024.
637
- 638 Cecilia Ying and Stephen Thomas. Label errors in BANKING77. In Shabnam Tafreshi, João Sedoc,
639 Anna Rogers, Aleksandr Drozd, Anna Rumshisky, and Arjun Akula (eds.), *Proceedings of the
640 Third Workshop on Insights from Negative Results in NLP*, pp. 139–143, Dublin, Ireland, May
641 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.insights-1.19. URL
642 <https://aclanthology.org/2022.insights-1.19/>.
643
- 644 Xiang Zhang, Junbo Jake Zhao, and Yann LeCun. Character-level convolutional networks for text
645 classification. In *NIPS*, 2015.
646
- 647 Pei Zhou, Andrew Zhu, Jennifer Hu, Jay Pujara, Xiang Ren, Chris Callison-Burch, Yejin Choi, and
648 Prithviraj Ammanabrolu. I cast detect thoughts: Learning to converse and guide with intents and
649 theory-of-mind in dungeons and dragons, 2023. URL <https://arxiv.org/abs/2212.10060>.

648 Yang Zhou, Hongyi Liu, Zhuoming Chen, Yuandong Tian, and Beidi Chen. Gsm-infinite: How do
649 your llms behave over infinitely increasing context length and reasoning complexity?, 2025. URL
650 <https://arxiv.org/abs/2502.05252>.

651 Andrew Zhu, Karmanya Aggarwal, Alexander Feng, Lara J. Martin, and Chris Callison-Burch.
652 FIREBALL: A dataset of dungeons and dragons actual-play with structured game state infor-
653 mation. In Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (eds.), *Proceedings of the*
654 *61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Pa-*
655 *pers)*, pp. 4171–4193, Toronto, Canada, July 2023a. Association for Computational Linguis-
656 tics. doi: 10.18653/v1/2023.acl-long.229. URL [https://aclanthology.org/2023.](https://aclanthology.org/2023.acl-long.229/)
657 *acl-long.229/*.

658 Andrew Zhu, Lara Martin, Andrew Head, and Chris Callison-Burch. Calypso: Llms as dungeon
659 master’s assistants. *Proceedings of the AAAI Conference on Artificial Intelligence and Interac-*
660 *tive Digital Entertainment*, 19(1):380–390, Oct. 2023b. doi: 10.1609/aiide.v19i1.27534. URL
661 <https://ojs.aaai.org/index.php/AIIDE/article/view/27534>.

662

663

664

665

666

667

668

669

670

671

672

673

674

675

676

677

678

679

680

681

682

683

684

685

686

687

688

689

690

691

692

693

694

695

696

697

698

699

700

701

702 A ICL LABEL VALIDATION
703

704 For each ICL dataset used to construct OOLONG-synth, we provide example instances that failed
705 validation (and thus were discarded) and that passed validation and were used for constructing con-
706 text windows. These examples are selected nearly-randomly; many of the instances that fail valida-
707 tion contain sexual content or offensive language, and we screen these out of the examples shown
708 wherever possible. While the instances screened out vary by dataset, we note that the datasets that
709 use labels scraped from web content (i.e. Yahoo Topics, AGNews, IMDB reviews, App Reviews)
710 generally have higher rates of validation failures. Some of these failures appear to be genuine mis-
711 labelings, which is likely because of noise in the user behavior (e.g. a user asks a question about
712 Business & Finance but mistakenly posts it in the Sports topic, or a user writes a review with mostly
713 negative text but gives the product 4 stars).

714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755

Input	Label	Passed validation?
‘Virtual Girlfriend’ Demands Gifts HONG KONG - She needs to be coddled with sweet talk and pampered with gifts, but you’ll never see her in the flesh. A Hong Kong company has developed a “virtual girlfriend” for new cell phones with video capability...	World	✗
Court Hears Case of Brain Damaged Woman (AP) AP - The Florida Supreme Court questioned lawyers Tuesday about the extent of the power handed to Gov. Jeb Bush under a law that let him order the reinsertion of a brain-damaged woman’s feeding tube.	Sci/Tech	✗
Rooney backs fiercesome threesome (AFP) AFP - Wayne Rooney believes the three-pronged attack of himself, Michael Owen and Jermain Defoe can put England on the fast track to the World Cup finals.	World	✗
From mouths of babes: What’s hot, what’s not How hot are the “Hot Dozen” toys? To find out, The Boston Globe put six toys from the 2004 Toy Wishes magazine Hot Dozen list in front of 13 kids from the Charlestown Boys and Girls Club.	Business	✗
Crematory operator to get 12 years Ray Brent Marsh, who is to enter the plea Friday, had faced up to 8,000 years in a case that shocked the nation two years ago when investigators found hundreds of corpses at his rural northwest facility.	Business	✗
Vending Machines Making Room for Healthy Products WASHINGTON (AP) – The typical vending machine fare consists of chocolate bars and potato chips, leaving few options for people seeking low-calorie or low-salt snacks. That is changing now as companies develop markets for products they expect to satisfy both nutritionists and consumers...	Sci/Tech	✓
Venezuelans Line Up to Vote on Chavez CARACAS, Venezuela - Summoned by bugle calls and the detonations of huge firecrackers, Venezuelans turned out in unprecedented numbers Sunday to vote on whether to force leftist President Hugo Chavez from office. Some lines at polling places extended for 1.25 miles...	World	✓
Singh Snares PGA Title Vijay Singh outlasts Justin Leonard and Chris DiMarco in a three-way playoff to win the PGA Championship on Sunday at Whistling Straits in Haven, Wisconsin.	Sports	✓
China’s Panchen Lama visits Tibet The boy named by the Chinese authorities as the 11th Panchen Lama visits a temple in Tibet.	World	✓
The Region’s Highest-Paid Executives Pay for the Washington area’s top executives rose significantly last year, reversing the downward trend that set in with the recession in 2001.	Business	✓

Table 5: AGNews examples that failed and passed validation.

Input	Label	Passed validation?
Kill Me..	positive	✗
Used to be great app. But since the last 7 updates always get - Authentication failed for: x@gmail.com	positive	✗
App force closing on changing THEME PLZ FIX!	positive	✗
You can't uninstall & I don't like my space being used. But am glad for the blind & deaf	positive	✗
GOOGLE TALKBACK Read carefully & then give at least a good reviews this app is made only for persons who got disability such as a blind person & if you want to get rid this simply just ROOT your device dumb head!!	negative	✗
This thing is great. I am at the bottom of the learning curve but in the few minutes I have played with it it seems like it will be easy to learn and to use. Just what I hoped for.	positive	✓
Crash when set a demodulation! :(negative	✓
doesn't work with rtl-2832 Says it works with rtl devices but doesn't. never displays anything even though the rtl 2832 works with other spectrum analyzers. no help files anywhere so good luck figuring out what's wrong!	negative	✓
Great app! Works really well with Galaxy note 5	positive	✓
Works great I am using HackRF One with this apps.	positive	✓

Table 6: App Reviews examples that failed and passed validation.

B QUESTION TYPES

B.1 OOLONG-SYNTH

Table 15, Table 16, and Table 17 list the questions used for each type of task in OOLONG-synth.

B.2 OOLONG-REAL

Table 18 and Table 19 lists the question types used for single and multi-episode settings in OOLONG-real dataset.

B.3 FULL RESULTS

In the tables below, we present the full results across both OOLONG-synth and OOLONG-real.

C DEEPSEEK R1 TRACE EXAMPLES

Figure 5 and Figure 6 contrast the behavior of this model on representative examples from the OOLONG-synth and OOLONG-real splits. See the main text for more discussion.

	Input	Label	Passed validation?
864			
865			
866			
867			
868			
869			
870			
871			
872			
873			
874			
875			
876			
877			
878			
879	The Unit secret agent has signed on to appear in three episodes as a businessman being shown some houses - and, presumably, one bedroom in particular - by our randy realtor heroine.	formal	✗
880			
881			
882			
883	Barely 12, with large brown eyes and stick-like arms, Fandi is 3 years older than his brother – in his eyes almost a man.	formal	✗
884			
885	D-Lister Avril Lavigne appears on the cover of Z-List magazine Savvy this month.	formal	✗
886			
887	ORHS is going through it now, but from all I can perceive the district is out for the quick fix, chop off the principal, instead of really trying to assist the students and dig to the bottom to find the truth.	formal	✗
888			
889			
890			
891	Little Sally Draper (Kiernan Shipka) is a "Patty Duke"-era girl - indeed, "Mad Men" is currently set during the month "The Patty Duke Show" premiered in 1963 - and on a "Patty Duke"-like show, she'd just be a lispy Shirley Temple doll with a crush on Daddy.	formal	✗
892			
893			
894			
895			
896	I have tried everything possible to attract business.	formal	✓
897			
898	Never heard that one before.	informal	✓
899			
900	"Yup."	informal	✓
901			
902	That's how we pass our traits to the next generation: through DNA	formal	✓
903			
904	Wedding rituals differ in different regions and communities in India.	formal	✓
905			
906			
907			
908			
909			
910			
911			
912			
913			
914			
915			
916			
917			

Table 7: Formality examples that failed and passed validation.

918 919 920	Input	Label	Passed validation?
921 922 923 924 925 926	Good movie, very 70s, you can not expect much from a film like this. Sirpa Lane is an actress of erotic films, a nice body but nothing exceptional. Not demand a lot from these films are light years away from the movies today, the world has changed incredibly. The plot is simple and the actors not extraordinary.	positive	✗
927 928 929 930 931	This film Evil Breed: The legend of samhain contains very little thought or effort. It is ridiculed with specs of ultra fast "slasher" style death and plain disgusting acts of death. The acting was rated a D as the actors show very little ability, and the stupidity of them in the film is too questionable.	positive	✗
932 933 934 935 936 937	The movie "Holly" may make the audience want to donate money towards organizations that improve the life for these poor youngsters, but the film's dramatic weaknesses may reduce its chances of being seen by enough people to make a difference. Overall, I think the concept is better as a documentary and it was not as touching as a movie.	positive	✗
938 939 940 941 942	Although this film put Davis on the map due to her brilliantly intense performance, this film is strangely unsatisfying to me as a whole. What I cannot fathom for the life of me is just how or why Phillip would take the constant abuse this tramp constantly dishes out towards him.	positive	✗
943 944 945 946 947	This film is just plain horrible. John Ritter doing pratt falls, 75% of the actors delivering their lines as if they were reading them from cue cards, poor editing, horrible sound mixing, and a plot that really goes nowhere. If I could sum this film up in one word, that word would be: Suckotrocity	negative	✓
948 949 950 951 952 953 954 955 956 957 958	Zentropa has much in common with The Third Man, another noir-like film set among the rubble of postwar Europe. Like TTM, there is much inventive camera work. There is an innocent American who gets emotionally involved with a woman he doesn't really understand, and whose naivety is all the more striking in contrast with the natives. But I'd have to say that The Third Man has a more well-crafted storyline. Zentropa is a bit disjointed in this respect. Perhaps this is intentional: it is presented as a dream/nightmare, and making it too coherent would spoil the effect. This movie is unrelentingly grim—"noir" in more than one sense; one never sees the sun shine. Grim, but intriguing, and frightening.	positive	✓
959 960 961 962 963	Never in my life have I come across a movie as bad as The Zombie Chronicles. Filmed on a budget of what looks to be about 20 bucks, TZC is a completely horrible horror movie that relies on lame, forgettable actors. Simply put, avoid TZC like a sexually-transmitted disease.	negative	✓
964 965 966 967 968 969	Without wishing to be a killjoy, Brad Sykes is responsible for at least two of the most dull and clichéd films I've ever seen. The acting is terrible, the print is shoddy, and everything about this film screams "seriously, you could do better yourself".	negative	✓

Table 8: IMDB examples that failed and passed validation. Only eight examples are shown because of the example length.

972

973

974

975

976

977

978

979

980

981

982

983

984

985

986

987

988

989

990

991

992

993

Input	Label	Passed validation?
Draupadi's eyes were diamonds. \leftrightarrow Draupadi's eyes were beautiful.	correct	✓
The faculty meeting was an easy breeze \leftrightarrow The faculty meeting was very easy and relaxing	correct	✓
The house was a tomb. \leftrightarrow The house was big.	incorrect	✓
I had already planted the idea in her mind. \leftrightarrow I had already scared her about the idea.	incorrect	✓
It is sad to observe the fruits of ignorance. \leftrightarrow It is sad to observe the effects of ignorance.	correct	✓

Table 9: Metaphors examples that passed validation. No examples failed validation for this dataset.

988

989

990

991

Input	Label	Passed validation?
You're safe. \rightarrow You have nothing to worry about.	neutral	✗
asks Burton, cupping his ear. \rightarrow Burton wanted to ask what was going on, but he couldn't do much of anything while bound and gagged.	contradiction	✗
I "poison my dearest Emily!" \rightarrow How could you think I would poison Emily?	entailment	✗
When he finally succeeded, after a prolonged siege and heavy losses, he punished the local population by cutting off the noses and lips of all men except those who played wind instruments. \rightarrow All men who played wind instruments were tasked with helping cut off people's noses.	neutral	✗
They died slowly, their eyes bulging and faces turning blue. \rightarrow They died asphyxiated.	neutral	✗
well you see that on television also \rightarrow You can see that on television, as well.	entailment	✓
Vrenna and I both fought him and he nearly took us. \rightarrow Neither Vrenna nor myself have ever fought him.	contradiction	✓
This analysis pooled estimates from these two studies to develop a C-R function linking PM to chronic bronchitis. \rightarrow The analysis proves that there is no link between PM and bronchitis.	contradiction	✓
He turned and smiled at Vrenna. \rightarrow He smiled at Vrenna who was walking slowly behind him with her mother.	neutral	✓
We sought to identify practices that were commonly implemented by the agencies within the past 5 years. \rightarrow We want to identify practices commonly used by agencies in the last 5 years	entailment	✓

1023

1024

1025

Table 10: MultiNLI examples that failed and passed validation.

1026

1027

1028

1029

1030

1031

1032

1033

1034

1035

1036

1037

1038

1039

1040

1041

1042

1043

1044

1045

1046

1047

1048

1049

1050

1051

1052

1053

1054

1055

1056

1057

1058

1059

1060

1061

1062

1063

1064

1065

1066

1067

1068

1069

1070

1071

1072

1073

1074

1075

1076

1077

1078

1079

Input	Label	Passed validation?
No care is anything that serves as an enticement.	True	✗
No care may be anything that serves as an enticement.	True	✗
A loud utterance is an appropriate definition of produce in no context.	True	✗
Complete attention is an appropriate definition of candidate in no context.	True	✗
No stress is a message received and understood.	True	✗
Action refers to a military engagement.	True	✓
Action never stands for a distinguishing quality.	True	✓
Action does not stand for a military engagement.	False	✓
Action may stand for a military engagement.	True	✓
Not a single action is a military engagement.	False	✓

Table 11: Negation examples that failed and passed validation.

Input	Label	Passed validation?
Go chase after her and run her over while she's crossing the street	ham	✗
i want to grasp your pretty booty :)	ham	✗
No da if you run that it activate the full version da.	ham	✗
i am seeking a lady in the street and a freak in the sheets. Is that you?	ham	✗
Should i send you naughty pix? :)	ham	✗
WINNER!! As a valued network customer you have been selected to receive a £900 prize reward! To claim call 09061701461. Claim code KL341. Valid 12 hours only.	spam	✓
Had your mobile 11 months or more? U R entitled to Update to the latest colour mobiles with camera for Free! Call The Mobile Update Co FREE on 08002986030	spam	✓
I'm gonna be home soon and i don't want to talk about this stuff anymore tonight, k? I've cried enough today.	ham	✓
SIX chances to win CASH! From 100 to 20,000 pounds txt CSH11 and send to 87575. Cost 150p/day, 6days, 16+ TsandCs apply Reply HL 4 info	spam	✓
I HAVE A DATE ON SUNDAY WITH WILL!!	ham	✓

Table 12: Spam examples that failed and passed validation. The examples that failed validation were almost exclusively ham messages with sexual content; the selected messages are the least inappropriate of the examples that failed validation.

1080

1081

1082

1083

1084

1085

1086

1087

1088

1089

1090

1091

1092

1093

1094

1095

1096

Table 13: TREC-coarse-Q examples that failed and passed validation. Only two examples failed validation.

Input	Label	Passed validation?
What explosive do you get by mixing charcoal, sulfur and saltpeter?	entity	✗
What is the procedure called for drilling a hole in your skull to achieve a higher consciousness?	entity	✗
Name 11 famous martyrs.	human being	✓
What's the Olympic motto?	description and abstract concept	✓
What is the highest waterfall in the United States?	location	✓
What does the abbreviation AIDS stand for?	abbreviation	✓
How many points make up a perfect fivepin bowling score?	numeric value	✓

1097

1098

1099

1100

1101

1102

1103

1104

1105

1106

1107

1108

1109

1110

1111

1112

1113

1114

1115

1116

1117

1118

1119

1120

1121

1122

1123

1124

1125

1126

1127

1128

1129

1130

1131

1132

1133

Input	Label	Passed validation?
Is there a God? The question to end all questions, and begin them.	Business & Finance	✗
why do we need to lie?	Health	✗
try to type the word supercalifragilisticexpialidocious. 20 times fast you can only make 10 mistakes? This is for fun!	Education & Reference	✗
do you think it is okay to tell a lie?	Education & Reference	✗
where can you purchase cesium carbonate?	Science & Mathematics	✗
What are good sources to find out about new gospel artists? Is there a site that focuses primarily on gospel?	Entertainment & Music	✓
How a black hole is formed? I would like to know how a black hole can possibly be formed. Are there any experimental evidence of such creation?	Science & Mathematics	✓
Economics of running a restaurant? Running a restaurant looks like hard work and long hours. What percentage of restaurants are profitable?	Business & Finance	✓
Why doesn't the NBA implement a minor leagues? I don't want to see any more High School kids on the court.	Sports	✓
how does a helicopter fly this is a miracle- I always wanted to learn to fly one of those. Can someone explain how can it get airborne?	Science & Mathematics	✓

Table 14: Yahoo examples that failed and passed validation. One class of examples that failed validation, not demonstrated here, are sexually explicit questions labeled as seemingly random categories (likely from spam posters on the original Yahoo Answers forum).

1134
1135
1136
1137

1138 Counting

- 1139 In the above data, which of the labels is the most common? Give your final answer in the form
1140 ‘label: answer’ where answer is one of the labels: {label_list}.
- 1141 In the above data, which of the labels is the least common? Give your final answer in the form
1142 ‘label: answer’ where answer is one of the labels: {label_list}.
- 1143 In the above data, is label ‘{A}’ more common, less common, or the same frequency as label
1144 ‘{B}’? Give your final answer in the form ‘Answer: {A} is [X]{B}’, where [X]is ‘more
1145 common than’, ‘less common than’, or ‘same frequency as’.
- 1146 In the above data, how many data points should be classified as label ‘{label}’? Give your final
1147 answer in the form ‘Answer: number’.
-

1149

1150 Table 15: Counting questions used in the OOLONG-synth dataset.

1151
1152
1153
1154
1155
1156
1157
1158

1159
1160 User

- 1161 In the above data, which user is represented most often? Give your final answer in the form
1162 ‘User: [X]’, where [X]is the user ID.
- 1163 In the above data, which user is represented the second most often? Give your final answer in
1164 the form ‘User: [X]’, where [X]is the user ID.
- 1165 For the following question, only consider the subset of users with IDs {user_names}. Among
1166 these users, which user is represented most often? Give your final answer in the form ‘User:
1167 [X]’, where [X]is the user ID.
- 1168 For the following question, only consider the subset of users with IDs {user_names}. Among
1169 these users, which user is represented the second most often? Give your final answer in the
1170 form ‘User: [X]’, where [X]is the user ID.
- 1171 For the following question, only consider the subset of instances that are associated with user
1172 IDs {user_names}. Among instances associated with these users, {any of the Counting ques-
1173 tions above}
- 1174 For the following question, only consider the subset of users with IDs {user_names}. Among
1175 these users, which user has the most instances with the label {label}? Give your final answer
1176 in the form ‘User: [X]’, where [X]is the user ID.
- 1177 In the above data, which user has the most instances with the label {label}? Give your final
1178 answer in the form ‘User: [X]’, where [X]is the user ID.
- 1179 In the above data, which user has more instances with the label {label}: User {A} or User {B}?
1180 Give your final answer in the form ‘User: [X]’, where [X]is the user ID.
-

1183
1184
1185
1186
1187

1188 Table 16: User questions used in the OOLONG-synth dataset.

1188
 1189
 1190
 1191
 1192
 1193
 1194
 1195
 1196
 1197
 1198
 1199

Timeline

1200 In the above data, which date is represented most often? Give your final answer in the form
 1201 ‘Date: [X]’, where [X]is the date in the format MM/DD/YYYY.
 1202 In the above data, which date is represented most often? Give your final answer in the form
 1203 ‘Date: [X]’, where [X]is the date in the format MM/DD/YYYY.
 1204 In the above data, which date is represented the second most often? Give your final answer in
 1205 the form ‘Date: [X]’, where [X]is the date in the format MM/DD/YYYY.
 1206 In the above data, how many dates are represented exactly {n} times? Give your final answer
 1207 in the form ‘Answer: [X]’, where [X]is the number of dates represented exactly {n} times.
 1208 In the above data, was label ‘{key}’ more common, less common, or the same frequency before
 1209 {time}, as compared to after {time}? Give your final answer in the form ‘Answer: {key} is
 1210 [X]before {time}’, where [X]is ‘more common’, ‘less common’, or ‘the same frequency’.
 1211 In the above data, was label ‘{key}’ more common, less common, or the same frequency before
 1212 {time}, as compared to after {time}? Give your final answer in the form ‘Answer: {key} is
 1213 [X]before {time}’, where [X]is ‘more common’, ‘less common’, or ‘the same frequency’.
 1214 For the following question, only consider the subset of instances that occur in {month_name} of
 1215 any year. Among instances occurring in {month_name},{any of the Counting questions above}
 1216 For the following question, only consider the subset of instances that occur between
 1217 {starting_date} and {ending_date}, inclusive. Among instances occurring in this date range,
 1218 {any of the Counting questions above}
 1219 In which month did the label ‘{label1}’ first occur more often than the label ‘{label2}’? Give
 1220 your final answer in the form ‘Answer: [month][year]’, where [month]is the name of the month
 1221 and [year]is the four-digit year where ‘{label1}’ first occurred more often than ‘{label2}’.
 1222 For how many months does the label ‘{label1}’ occur more frequently than the label
 1223 ‘{label2}’? Disregard months where there is a tie. Give your final answer in the form ‘Answer:
 1224 [X]’, where [X]is the number of months where ‘{label1}’ occurs more often than ‘{label2}’.
 1225 For how many months is the label ‘{label}’ the single most frequently occurring label? Disre-
 1226 gard months where there is a tie for the most common label. Give your final answer in the form
 1227 ‘Answer: [X]’, where [X]is the number of months where ‘{label}’ is the most common label.

1231
 1232
 1233
 1234
 1235
 1236
 1237
 1238
 1239
 1240
 1241

Table 17: Timeline questions used in the OOLONG-synth dataset.

1242	
1243	
1244	
1245	
1246	
1247	
1248	
1249	
1250	
1251	Rolls
1252	Total number of rolls in this episode?
1253	Total number of rolls by the character {character name} in this episode?
1254	Total number of rolls by the player {player name} in this episode?
1255	Total number of rolls of type {roll type} in this episode?
1256	Number of rolls of natural value {roll value} in this episode?
1257	In this episode, what percentage of rolls were of value {roll value}? round to the nearest integer.
1258	What is the most common roll type in this episode? Return a comma separated list.
1259	What is the least common roll type in this episode? Only include types with more than one roll.
1260	Return a comma separated list.
1261	What is the most common natural roll value in this episode? Return a comma separated list.
1262	What is the least common natural roll value in this episode? Only include values with more
1263	than one roll. Return a comma separated list.
1264	What is the count of Crits? (natural rolls of value 1 or 20)?
1265	What is the count of Nat20s (natural rolls of value 20)?
1266	What is the count of Nat1s (natural rolls of value 1)?
1267	Spells
1268	How many spells were cast during this episode?
1269	How many spells were cast by the character {character name} in this episode?
1270	How many spells were cast by the player {player name} in this episode?
1271	How many {spell type} spells were cast during this episode?
1272	What are the first {count} spells cast in this episode? Return a comma separated list.
1273	What are the last {count} spells cast in this episode? Return a comma separated list.
1274	Return a comma separated list and retain the order of spells as they appear in the episode.
1275	What is the last spell cast by each character in this episode? Return a comma separated list and
1276	retain the order of spells as they appear in the episode.
1277	How many characters cast {spell name} spell in this episode?
1278	What is the most common spell in this episode? Return a comma separated list.
1279	What is the least common spell in this episode? Only include spells that were cast at least once.
1280	Return a comma separated list.
1281	Which spells were cast by more than one character in this episode? Return a comma separated
1282	list.
1283	What is the total number of cantrip spells cast in this episode?
1284	In this episode, how many times was a spell cast at a level higher than its base level?
1285	In this episode, which spells were cast at a level higher than their base level? Return a comma
1286	separated list of unique spells.
1287	Table 18: Question types used in the OOLONG-real dataset (single episode).
1288	
1289	
1290	
1291	
1292	
1293	
1294	
1295	

1296

1297

1298

1299

1300

1301

Rolls

1302

What is the cummulative total of rolls by the end of episode {episode index}? Count the number of rolls and not the values of the rolls.

1304

1305

What is the cummulative total of rolls by the character {character name} at the end of episode {episode index}? Count the number of rolls and not the values of the rolls.

1306

Total number of rolls across all the episodes?

1307

Total number of rolls by the character {character name} across all episodes?

1308

Total number of rolls by the player {player name} across all episodes?

1309

Total number of rolls of type {roll type} across all episodes?

1310

Number of rolls of natural value {roll value} across all episodes?

1311

Across all episodes, what percentage of rolls were of value {roll value}? round to the nearest integer.

1312

What is the most common roll type across all episodes? Return a comma separated list.

1313

What is the least common roll type across all episodes? Only include types with more than one roll. Return a comma separated list.

1314

What is the most common natural roll value across all episodes? Return a comma separated list.

1315

What is the least common natural roll value across all episodes? Only include values with more than one roll. Return a comma separated list.

1316

What is the total count of Crits across all episodes? (natural rolls of value 1 or 20)?

1317

What is the total count of Nat20s across all episodes? (natural rolls of value 20)?

1318

What is the total count of Nat1s across all episodes? (natural rolls of value 1)?

1322

Spells

1323

What is the cummulative total of spells cast by the end of episode {episode index}?

1324

What is the first spell cast in the episode {episode index}?

1325

What is the second spell cast in the episode {episode index}?

1326

What is the third spell cast in the episode {episode index}?

1327

List the first spell cast in each episode? Return a comma separated list.

1328

List the last spell cast in each episode? Return a comma separated list.

1329

List the first spell cast by the character {character name} in each episode? Return a comma separated list.

1330

List the last spell cast by the character {character name} in each episode? Return a comma separated list.

1331

How many spells were cast across all episodes?

1332

How many spells were cast by the character {character name} across all episodes?

1333

How many spells were cast by the player {player name} across all episodes?

1334

How many {spell type} spells were cast across all episodes?

1335

How many characters cast {spell name} spell across all episodes?

1336

What is the most common spell across all episodes? Return a comma separated list.

1337

What is the least common spell across all episodes? Only include spells that were cast at least once. Return a comma separated list.

1338

What is the total number of cantrip spells cast across all episodes?

1339

Across all episodes, how many times was a spell cast at a level higher than its base level?

1340

Across all episodes, which spells were cast at a level higher than their base level? Return a comma separated list of unique spells.

1344

Table 19: Question types used in the OOLONG-real dataset (multi-episode).

1345

1346

1347

1348

1349

1350

1351

1352

1353

1354

1355

1356

1357

1358

1359

1360

1361

1362

1363

1364

1365

1366

1367

1368

1369

1370

1371

1372

1373

1374

1375

1376

1377

1378

1379

1380

1381

1382

1383

1384

1385

1386

1387

1388

1389

1390

1391

1392

1393

1394

1395

1396

1397

1398

1399

1400

1401

1402

1403

<think> We are given 21 pairs. We need to classify each pair as one of ‘neutral’, ‘contradiction’, or ‘entailment’. Then, we will count the frequency of ‘contradiction’ and ‘entailment’ and compare them. Let’s go through each pair one by one: 1. “and then the class uh double A Dallas Fort Worth Spurs we later uh they went to the Texas league - Dallas Fort Worth Spurs got into the Texas League.” - The first sentence says they went to the Texas league. The second says they got into the Texas League. This seems to be the same meaning. So, entailment. 2. “Specifically, although most public companies are required to have a board that is comprised of a majority of independent directors and certain key committees are required to be comprised solely of outside directors, a closer look reveals that the independence of many boards may be more a matter of form than substance. - Public companies can put whoever they want on their boards.” - The first sentence says that public companies are required to have independent directors, but the independence might be in form only. The second sentence says they can put whoever they want. This contradicts the requirement mentioned. So, contradiction. **[....authors truncated classifications for space]** 19. “And the bureaucratic approach prevents the section as a whole from taking positions on—or creating—literary issues. - Bureaucracy makes groups not take positions or a stand on literary issues.” - The second is a restatement of the first. Entailment. 20. “They are linked not only by geographical location but also by here are the upmarket and fashionable neighborhoods, with affluent residential sections, trendy restaurants, and L.A.’s fabulous shopping areas. - They are linked because they both have very exclusive shoe stores.” - The first says they are linked by upmarket and fashionable neighborhoods, etc. The second says they are linked because of exclusive shoe stores. The first does not mention shoe stores. So, the second is a specific reason that is not stated. This could be neutral. However, note that the first lists shopping areas, which might include shoe stores, but it doesn’t specify. The second is too

Figure 5: Example reasoning trace for OOLONG-synth. The model fails to complete the task because it runs out of tokens during the reasoning trace. The majority of the trace (classifying items 3-18) is omitted for space.

1404

1405

1406

1407

1408

1409

1410

1411

1412

1413

1414

1415

1416

1417

1418

1419

1420

1421

1422

1423

1424

1425

1426

1427

1428

1429

1430

1431

1432

1433

1434

1435

1436

1437

1438

1439

1440

1441

1442

1443

1444

1445

1446

1447

1448

1449

1450

1451

1452

1453

1454

1455

1456

1457

Based on the provided transcript of Episode 2, the second spell cast is Disintegrate, cast by Lyra (Felicia Day's character).

Sequence of Events: 1. First Spell: Vex'ahlia (Laura Bailey) casts Hunter's Mark on a frost giant (during her turn in combat). 2. Second Spell: Lyra (Felicia Day) casts Disintegrate on the same frost giant (during her subsequent turn).

Key Excerpt: Felicia: "I have a spell called Disintegrate. Can I cast that on its wing? [...] I'm doing Disintegrate. And it's 10d6 plus a d10, plus 40." Matt: "Disintegrate is nasty. [...] As you release the torrent of electrical energy, it courses through the body [...] the giant looks severely messed up."

Thus, the second spell in Episode 2 is Disintegrate.

Disintegrate

Figure 6: Example reasoning trace for OOLONG-real. The model completes the task within a relatively short reasoning trace, despite this task using substantially longer input.